

RAG

RAG = Retrieval Augmented Generation

Retrieval-Augmented Generation (RAG) is the process of optimizing the output of a large language model, so it references an authoritative knowledge base outside of its training data sources before generating a response. Large Language Models (LLMs) are trained on vast volumes of data and use billions of parameters to generate original output for tasks like answering questions, translating languages, and completing sentences. RAG extends the already powerful capabilities of LLMs to specific domains or an organization's internal knowledge base, all without the need to retrain the model. It is a cost-effective approach to improving LLM output so it remains relevant, accurate, and useful in various contexts.



WHY DO I NEED RAG?

First: what an LLM actually is

An LLM is:

- A probabilistic text generator trained on past data

Key properties:

- No live access to your data
- No memory of your files
- Fixed knowledge cutoff
- Limited context window
- Will confidently guess if unsure

So the real question becomes:

- How do we make an LLM answer using *your* data, correctly?

That is why RAG exists.

Problem 1: Knowledge cutoff

Situation

You ask:

"What does *my* company policy say about remote work?"

The LLM:

- Was trained on public internet data
- Has **never seen your policy**

Without RAG

- It guesses
- Sounds confident
- Is probably wrong ✗

With RAG

- Retrieve the actual policy
- Inject it into the prompt
- Answer grounded in facts ✓

↗️ RAG gives LLM access to private & new knowledge

Problem 2: Context window limits

LLMs can only see:

- A few thousand tokens (even large models have limits)

Situation

You upload:

- 300-page PDF
- Large codebase
- Research corpus

Without RAG

- You can't fit it all
- Earlier parts are forgotten
- Answers become inconsistent ✗

With RAG

- Store everything externally
- Retrieve only what's relevant
- Feed **just-in-time context** ✓

↗️ RAG solves scale

Problem 3: Hallucinations (the big one)

LLMs must answer — even when they don't know.

Situation

You ask:

"What experiment did the paper run in section 5?"

Without RAG

- The model may:
 - Invent an experiment
 - Mix sections
 - Fabricate details ✗

With RAG

- If section 5 exists → retrieved
- If not → nothing retrieved → "not found" response
- Grounded, verifiable answer ✓

↗ RAG constrains the LLM to reality

Problem 5: Explainability & trust

In real systems:

- Enterprises
- Research
- Legal
- Healthcare

You must answer:

"Where did this answer come from?"

Without RAG

- "The model says so" ✗

With RAG

- "This answer comes from:
 - Page 12
 - Section 3.1
 - Paragraph 2"

↗ RAG enables citations & auditability

Why “updating knowledge” is hard for LLMs

An LLM’s knowledge is stored **inside its weights**.

Think of the model like this:

kotlin

Copy

```
Training data —> billions of weight values —> behavior
```

Once trained:

- The model **does not have a database**
- It **cannot edit facts**
- It **cannot delete or insert knowledge cleanly**

So if knowledge **changes**, you have a serious problem.

Option 1: Fine-tuning (✗ why it's expensive & risky)

What fine-tuning actually means

Fine-tuning = changing the model's weights

java

```
old weights  
+ new training data  
+ gradient updates  
= new weights
```

This is not adding a file.

This is rewriting the brain.

Why fine-tuning is costly

1 Compute cost

- GPUs / TPUs
- Many epochs
- Memory + time
-

Even small fine-tunes can cost hundreds to thousands of dollars.

2 Data preparation cost

- Clean datasets
- Labeling
- Formatting
- Validation

Garbage data = broken model.

3 Iteration is slow

- Change policy?
- New research paper?
- Updated law?

➡ You must fine-tune again.

Catastrophic forgetting (critical concept)

When you fine-tune on new data:

pgsql

Old knowledge ← overwritten by → New knowledge

Example:

- Model knew general physics
- You fine-tune on company policy
- It starts answering physics questions worse ✗

This happens because:

- Gradients don't "respect boundaries"
- Updating weights affects everything
- ✗ You cannot safely isolate knowledge in weights

Why fine-tuning is bad for changing facts

Facts change often:

- Laws
- Policies
- Prices
- Research
- APIs

Fine-tuning is:

- ✗ Global
- ✗ Slow
- ✗ Risky
- ✗ Hard to undo

Why RAG is safer

1 No forgetting

- Model weights unchanged
- Core reasoning preserved

2 Isolation

- Each document is independent
- Bad doc ≠ broken model

3 Auditable

- You know exactly where knowledge lives
- Can trace answers to sources

Key mental model (very important)

Fine-tuning

"Let's change how the brain works."

RAG

"Let's give the brain a book to read before answering."

You don't retrain a doctor every time a medical guideline updates — you give them new documentation.

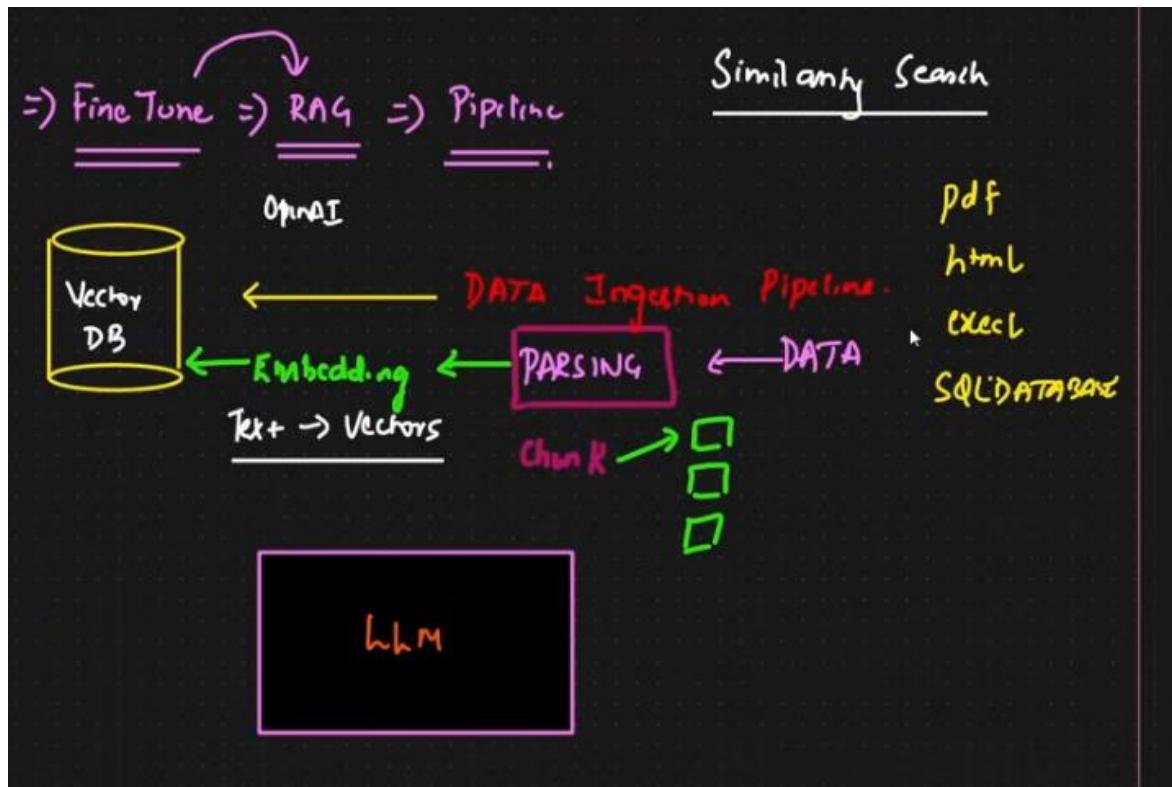


[Retrieval-Augmented Generation \(RAG\)](#) is a framework that combines the strengths of information retrieval and generative models:

- **Retriever:** The retriever component fetches relevant documents from a large corpus or knowledge base based on the input query.
- **Generator:** The generator then takes the retrieved documents and the query to generate a coherent and contextually relevant response.

It allows a model to retrieve relevant documents from a knowledge base and use those documents to augment the generation process, resulting in more accurate, context-aware and insightful responses. This approach has shown promising results in various applications such as question answering, dialogue systems and content generation. In this article we will build a RAG Application.

DATA INGESTION



Mental model

java

Copy code

```
Raw Data (PDFs, CSVs, APIs)
    ↓ ← ingestion
System Storage (DB, Data Lake, Vector DB)
    ↓ ← processing, parsing, embeddings
LLM / AI pipeline
```

Summary:

Data ingestion = feeding raw data into your system so it can be stored, processed, and used by AI or other software.

STEP -1

DATA PARSING:

Step 1: Data Parsing / Loading

In any RAG system, the very first step is:

Take your raw documents (PDF, TXT, etc.) and convert them into plain text the computer can work with.

This is called **data parsing**. Without it, embeddings or retrieval cannot happen.

```
1 load_txt(path)
python

def load_txt(path):
    with open(path, "r", encoding="utf-8", errors="ignore") as f:
        return f.read()
```

What it does:

1. Opens a `.txt` file at the given `path`.
2. Reads the entire content into a single string.
3. Uses `encoding="utf-8"` → handles international characters.
4. `errors="ignore"` → ignores weird/unreadable characters instead of crashing.

Why:

- TXT files are already text, but computers need **uniform strings** to process.
- This function ensures **robust reading** without errors.

```
2 load_pdf(path)
python

def load_pdf(path):
    reader = PdfReader(path)
    text = ""
    for p in reader.pages:
        t = p.extract_text()
        if t:
            text += t + "\n"
    return text
```

What it does:

1. Uses `PdfReader` from `pypdf` → reads PDF structure.
2. Loops over all **pages**.
3. Calls `extract_text()` → gets text content of the page.
4. Concatenates all pages into a **single string** separated by newline (`\n`).
5. Returns the **entire document as plain text**.

Why:

- PDFs are **complex files** (fonts, layout, images, tables).
- LLMs cannot read PDFs directly.
- We need a **clean, plain-text version** for embeddings and retrieval.

```
3 load_document(path)
python

def load_document(path):
    if path.lower().endswith(".txt"):
        return load_txt(path)
    elif path.lower().endswith(".pdf"):
        return load_pdf(path)
    return None

What it does:
1. Determines file type by extension (.txt or .pdf).
2. Calls the appropriate loader (load_txt or load_pdf).
3. Returns None if the file is unsupported.

Why:
• Makes the pipeline file-type agnostic.
• You can drop mixed file types into your docs folder and it works automatically.
```

Step 2: Chunking Process

After parsing, your documents are now **one long string of text**:

text

"This is a long document ... with hundreds or thousands of words..."

Problem: LLMs cannot process huge text in a single pass.

Chunking solves this by breaking text into **smaller, manageable pieces** (chunks).

Why Chunking is Needed

1. LLM context window limitation

- Every LLM can only see a **limited number of tokens** at once.
- Example: GPT-4 8k context → ~6,000 words max.
- Chunking ensures **each piece fits** in the model's context.

2. Semantic retrieval accuracy

- RAG retrieves **relevant chunks** from your document.
- Smaller, coherent chunks → **more precise retrieval**.
- Large chunks → LLM may miss the exact answer.

3. Maintain coherence

- Chunks should **preserve sentences / meaning**.
- Bad chunking → mid-sentence split → LLM may misinterpret.

4. Efficient embedding

- Embeddings are **vector representations** of chunks.
- If chunks are too big → embeddings **lose fine-grained detail**.
- If chunks are too small → retrieval is noisy (too many tiny pieces).

```
def chunk_text(text, size, strategy="fixed"):
    """
    Chunking strategies for experimental comparison

    Improved to maintain sentence boundaries and context
    """

    if strategy == "fixed":
        words = text.split()
        chunks = []

        for i in range(0, len(words), size):
            chunk = " ".join(words[i:i+size])
            # Clean up excessive whitespace
            chunk = " ".join(chunk.split())
            chunks.append(chunk)

    return chunks

    elif strategy == "sentence":
        # Sentence-based chunking (more coherent)
        try:
            import nltk
            sentences = nltk.sent_tokenize(text)
        except:
            # Fallback if NLTK not available
            sentences = text.split('. ')

        chunks = []
        current_chunk = []
        current_size = 0

        for sentence in sentences:
            sentence_words = len(sentence.split())

            if current_size + sentence_words > size and current_chunk:
                # Save current chunk and start new one
                chunks.append(" ".join(current_chunk))
                current_chunk = [sentence]
                current_size = sentence_words
            else:
                current_chunk.append(sentence)
                current_size += sentence_words

        # Add final chunk
        if current_chunk:
            chunks.append(" ".join(current_chunk))

    return chunks

# Default to fixed chunking
words = text.split()
return [" ".join(words[i:i+size]) for i in range(0, len(words), size)]
```

The `chunk_text` Function

python

```
def chunk_text(text, size, strategy="fixed"):  
    • text → string from parsed document  
    • size → number of words per chunk  
    • strategy → how to split:  
        • "fixed" → split every size words, ignoring sentences  
        • "sentence" → split at sentence boundaries, more coherent
```

1 Fixed Chunking

python

```
words = text.split()  
for i in range(0, len(words), size):  
    chunk = " ".join(words[i:i+size])  
    chunks.append(chunk)
```

- Simple word-based splitting
 - Easy and fast
 - Drawback: may split sentences mid-way → slightly incoherent
- Good for baseline experiments and testing speed vs accuracy.

2 Sentence-based Chunking

python

```
import nltk  
sentences = nltk.sent_tokenize(text)
```

- Tries to keep sentences together
- Accumulates sentences until reaching `size` limit
- Produces chunks that are semantically coherent

Fallback: If NLTK unavailable → split by '.'.

Why important:

- Coherent chunks → LLM understands context better → reduces hallucinations

How does NLTK know where to break sentences?

`nltk.sent_tokenize(text)` is a sentence tokenizer.

Its job is to split a paragraph or document into individual sentences.

1 Rule-based approach (Punkt tokenizer)

NLTK uses a pretrained sentence tokenizer called Punkt:

- It looks for sentence boundaries based on punctuation:
 - ., !, ? → usually end of a sentence
 - Newlines, capitalization patterns → clues for start of a new sentence
- Example:

python

 Copy code

```
text = "Hello world! How are you today? I hope you're fine."
sentences = nltk.sent_tokenize(text)
print(sentences)
```

Output:

css

 Copy code

```
['Hello world!', 'How are you today?', "I hope you're fine."]
```

2 Smart handling of abbreviations

- Punkt is trained on real text, so it avoids splitting in wrong places:

python

 Copy code

```
text = "Dr. Smith went to the hospital. He was late."
sentences = nltk.sent_tokenize(text)
print(sentences)
```

Output:

css

 Copy code

```
['Dr. Smith went to the hospital.', 'He was late.']}
```

- Notice it did not split at "Dr." — it knows abbreviations like "Dr.", "Mr.", "e.g.", "i.e." usually do not end sentences.

Summary Table			
Strategy	Size measures	Example chunking	Pros/Cons
fixed	words	5 words per chunk	Fast, predictable, may split sentences
sentence	sentences	2 sentences per chunk	Coherent, better for context, slightly slower

1 Fixed Chunking (`size` = words per chunk)

✓ Advantages

- Fast and simple
- Predictable chunk size → easy to tune embeddings and database
- Good for massive datasets where speed matters

✗ Disadvantages

- May split sentences mid-way → LLM sees incomplete ideas
- Risk of hallucinations or less coherent answers

⚡ Use Fixed When

- Doing experiments / benchmarks to test embeddings or retrieval speed
- Indexing huge volumes of text quickly
- Chunk sizes are small (20–50 words), so splitting sentences doesn't matter much
- You don't need perfect semantic coherence

2 Sentence-based Chunking (`size` = sentences per chunk)

✓ Advantages

- Preserves full sentences → coherent chunks
- LLM sees complete ideas, reducing hallucination
- Better for grounded answers, citations, QA

✗ Disadvantages

- Slightly slower (needs sentence tokenization, NLTK or similar)
- Chunk size may vary → sometimes fewer or more words than expected

⚡ Use Sentence-Based When

- Working with formal or academic documents (papers, reports)
- Answers must be high-quality, grounded, and coherent
- Chunk sizes are medium to large (>100 words)
- You need explainability: LLM can reference full sentences

Rule of Thumb

Goal / Scenario	Recommended Strategy
Fast indexing of huge document collections	Fixed
High-quality QA or summarization	Sentence
Small chunks (<50 words)	Fixed is fine
Medium/large chunks (>100 words)	Sentence preferred
Experiments / baseline tests	Fixed
Legal / academic / structured content	Sentence

Function: `load_and_chunk_folder`

```
python

def load_and_chunk_folder(folder, chunk_size, strategy="fixed"):
    chunks, metas = [], []

    for filename in os.listdir(folder):
        fp = os.path.join(folder, filename)
        text = load_document(fp)

        if not text:
            continue

        file_chunks = chunk_text(text, chunk_size, strategy)

        for i, c in enumerate(file_chunks):
            chunks.append(c)
            metas.append({"source": filename, "chunk": i})

    return chunks, metas
```

Step 1: Initialize lists

```
Python

chunks, metas = [], []
```

- `chunks` → will store all text chunks from all documents
- `metas` → will store metadata about each chunk (where it came from, index, etc.)

Metadata is important for retrieval traceability in RAG.

Step 2: Loop through all files in folder

```
Python

for filename in os.listdir(folder):
    fp = os.path.join(folder, filename)
```

- `os.listdir(folder)` → lists all files in the folder
- `fp` → full path to the file (`folder/filename`)

This allows the function to process multiple documents automatically.

Step 3: Load the document

```
Python

text = load_document(fp)
if not text:
    continue
```

- Uses `load_document` → parses TXT or PDF into plain text
- If parsing fails or file is unsupported → skip it

This ensures only valid text is processed.

Step 4: Chunk the text

python

Copy code

- ```
file_chunks = chunk_text(text, chunk_size, strategy)
```
- Uses the chunking function we discussed
  - `chunk_size` → number of words (fixed) or sentences (sentence strategy)
  - Produces a list of chunks for this document

#### Step 5: Store chunks and metadata

python

Copy code

```
for i, c in enumerate(file_chunks):
 chunks.append(c)
 metas.append({"source": filename, "chunk": i})
```

- Loops through each chunk `c` in `file_chunks`
  - Appends it to the master list of chunks
  - Appends metadata dictionary:
    - `"source"` → file name
    - `"chunk"` → chunk index in that file
- This is critical for RAG: when the LLM retrieves a chunk, you know which document and which part it came from.

#### Step 6: Return results

python

Copy code

```
return chunks, metas
```

- `chunks` → list of all text chunks from all files
- `metas` → corresponding metadata for traceability

These are ready for embedding and insertion into the vector database for retrieval.

#### 1. Load documents

- Convert PDFs, TXT, etc. → plain text
- Using `load_document` (previous step)

#### 2. Split text into manageable chunks

- LLMs cannot read huge documents all at once
- Chunks can be words-based (fixed) or sentence-based
- Ensures LLM can see full context in each chunk

#### 3. Attach metadata

- Keep track of source and chunk index
- Enables traceable, explainable answers

#### 4. Output for embeddings

- After this step, we have:
  - `chunks` → list of text pieces
  - `metas` → metadata about each piece
- These are ready to create embeddings, which will go into the vector database

#### Why this step is critical

- Without chunking: LLM might get truncated text → incomplete answers → hallucinations
- Without metadata: LLM cannot cite sources → answers become untrustworthy
- Without loading all documents: RAG has nothing to retrieve → no knowledge

#### Analogy

Imagine you're building a searchable library for an AI assistant:

1. You have books (documents)
2. You cut them into pages (chunks)
3. You label each page with book name and page number (metadata)
4. Now your AI can search pages quickly and answer questions grounded in the right sources

#### Summary

- Step goal: Convert raw documents → coherent chunks + metadata → ready for embedding & retrieval
- End result: `chunks` + `metas` that the RAG pipeline can store, search, and use for LLM answers

## Step 3: Embedding

```
EXPERIMENTAL_PIPELINES = {
 "miniLM_fixed100": {
 "model": "all-MiniLM-L6-v2",
 "chunk_size": 100,
 "chunk_strategy": "fixed",
 "description": "Fast, small embeddings with small fixed chunks"
 },
 "miniLM_fixed200": {
 "model": "all-MiniLM-L6-v2",
 "chunk_size": 200,
 "chunk_strategy": "fixed",
 "description": "Fast embeddings with larger context windows"
 },
 "E5small_fixed150": {
 "model": "intfloat/e5-small",
 "chunk_size": 150,
 "chunk_strategy": "fixed",
 "description": "E5 embeddings optimized for semantic search"
 },
 "BGE_fixed150": {
 "model": "BAAI/bge-base-en-v1.5",
 "chunk_size": 150,
 "chunk_strategy": "fixed",
 "description": "BGE embeddings with balanced performance"
 },
 "BGE_fixed250": {
 "model": "BAAI/bge-base-en-v1.5",
 "chunk_size": 250,
 "chunk_strategy": "fixed",
 "description": "BGE embeddings with larger context (better for
complex queries)"
 },
}

EMBEDDERS = {}
for name, cfg in EXPERIMENTAL_PIPELINES.items():
 print(f"[EXPERIMENT INIT] Loading embedding model:
{cfg['model']}")
 EMBEDDERS[name] = SentenceTransformer(cfg["model"])
```

## 1 The Experimental Pipelines

Your code defines five different pipelines, each specifying:

- `model` → which embedding model to use
- `chunk_size` → how big each chunk should be
- `chunk_strategy` → "fixed" or "sentence"
- `description` → notes on speed, accuracy, or intended use

This is used to experiment with RAG performance across different embedding approaches.

## 2 Embedding Models Explained

### a) `all-MiniLM-L6-v2` (used in `miniLM_fixed100` and `miniLM_fixed200`)

- Type: Small transformer model from SentenceTransformers
- Strengths:
  - Fast and lightweight → good for large-scale datasets
  - Handles general semantic similarity well
- Weaknesses:
  - Not as accurate on nuanced or technical text
  - Smaller embedding dimension → less context capture
- Best for:
  - Quick experiments
  - Short, simple documents (blogs, FAQs, short reports)

### b) `intfloat/e5-small` (used in `E5small_fixed150`)

- Type: E5 embeddings optimized for semantic search
- Strengths:
  - Good at capturing semantic meaning even if wording differs
  - Optimized for retrieval and question-answering
- Weaknesses:
  - Slightly slower than MiniLM
- Best for:
  - Medium-sized documents
  - Semantic search where user queries may be paraphrased

### c) `BAAI/bge-base-en-v1.5` (used in `BGE_fixed150` and `BGE_fixed250`)

- Type: BGE (Big Graph Embeddings) base model
- Strengths:
  - Balanced performance → good accuracy and speed
  - Handles complex or long documents well
  - Larger context window → better for detailed queries
- Weaknesses:
  - Slightly heavier → slower than MiniLM
- Best for:
  - Research papers, technical manuals, legal documents
  - Complex question-answering tasks

### 3 Chunk Size

- What it is: number of words per chunk (for `fixed`)
- Defined in the pipeline:

| Pipeline                      | Chunk Size | Notes                                           |
|-------------------------------|------------|-------------------------------------------------|
| <code>minilm_fixed100</code>  | 100 words  | Small chunks, fast retrieval                    |
| <code>minilm_fixed200</code>  | 200 words  | Larger chunks → more context                    |
| <code>E5small_fixed150</code> | 150 words  | Medium chunks for semantic search               |
| <code>BGE_fixed150</code>     | 150 words  | Balanced size for context vs speed              |
| <code>BGE_fixed250</code>     | 250 words  | Large chunks → more context for complex queries |

- Default: In your code, `chunk_size` is defined per pipeline, there is no "global default" for all.

### 4 Chunk Strategy

- What it is: how text is split into chunks

| Strategy              | How it works                                                             | When to use                                                   |
|-----------------------|--------------------------------------------------------------------------|---------------------------------------------------------------|
| <code>fixed</code>    | Split every <code>chunk_size</code> words                                | Fast indexing, simple documents                               |
| <code>sentence</code> | Split at sentence boundaries (~ <code>chunk_size</code> words per chunk) | Better coherence, reduces hallucinations, formal/complex text |

- In your pipelines: all are "`fixed`".
- Why: probably for speed and experimental consistency, but you could change it to "`sentence`" for academic or legal documents.

### 5 Choosing the Best Pipeline for Your Document

| Document Type                        | Recommended Pipeline          | Why                                            |
|--------------------------------------|-------------------------------|------------------------------------------------|
| Short FAQs, blogs                    | <code>minilm_fixed100</code>  | Fast, small chunks, good enough semantic match |
| Medium technical docs                | <code>E5small_fixed150</code> | Semantic search optimized, medium context      |
| Large, complex, or research papers   | <code>BGE_fixed250</code>     | Large chunks → more context, better grounding  |
| General documents, balanced accuracy | <code>BGE_fixed150</code>     | Fast enough with good accuracy                 |
| Quick experiments / benchmarks       | <code>minilm_fixed200</code>  | Large enough for context, fast to run          |

Key insight:

- Smaller models + small chunks → fast, less context
- Larger models + larger chunks → slower, more accurate, less hallucination

## 1 What is EMBEDDERS = {} ?

This is a dictionary that will store loaded embedding models.

Think of it as:

```
makefile

EMBEDDERS = {
 "minILM_fixed100": <MinILM embedding model>,
 "E5small_fixed150": <E5 embedding model>,
 "BGE_fixed250": <BGE embedding model>,
 ...
}
```

### ★ Why a dictionary?

Because later you want to say:

```
python

EMBEDDERS["BGE_fixed250"].encode(chunks)
```

So each experimental pipeline can use its own embedding brain.

## 2 What is EXPERIMENTAL\_PIPELINES.items() ?

EXPERIMENTAL\_PIPELINES looks like:

```
python

{
 "minILM_fixed100": {...},
 "minILM_fixed200": {...},
 "E5small_fixed150": {...},
 ...
}
```

Calling .items() gives:

```
python

(name, cfg)
```

Example:

```
python

name = "minILM_fixed100"
cfg = {
 "model": "all-MiniLM-L6-v2",
 "chunk_size": 100,
 "chunk_strategy": "fixed",
 "description": "Fast, small embeddings..."
}
```

So the loop runs once per pipeline configuration.

### 3 What does this line do?

```
python
```

```
SentenceTransformer(cfg["model"])
```

This is the most important part.

This line:

- ✓ Downloads the model (if not already downloaded)
- ✓ Loads it into memory (RAM / GPU)
- ✓ Prepares it to convert text → vector embeddings

For example:

```
python
```

```
SentenceTransformer("all-MiniLM-L6-v2")
```

loads a neural network that can do:

```
python
```

```
"text about neural networks"
+ [0.021, -0.34, 0.91, ..., 0.002] # vector
```

★ These vectors are what allow semantic search.

### 4 What does this line do?

```
python
```

```
EMBEDDERS[name] = SentenceTransformer(cfg["model"])
```

This stores the loaded model in the dictionary.

So after the loop finishes:

```
python
```

```
EMBEDDERS = {
 "minILM_fixed100": MiniLM model object,
 "minILM_fixed200": MiniLM model object,
 "E5small_fixed150": E5 model object,
 "BGE_fixed150": BGE model object,
 "BGE_fixed250": BGE model object
}
```

### 5 Why load models once here?

#### ✗ Bad approach (slow, wasteful)

```
python
```

```
for each chunk:
 model = SentenceTransformer("model-name")
 model.encode(chunk)
```

This would:

- Reload model thousands of times
- Kill performance
- Waste memory

#### ✓ Correct approach (what you did)

```
python
```

```
Load model ONCE
Reuse it for all chunks
```

★ This is production-grade design.

## VECTOR DATABASE:

```
=====
VECTOR DATABASE SETUP
=====

DBS = {}

for name in EXPERIMENTAL_PIPELINES.keys():
 db_path = f'./chroma_db_{name}'
 client = chromadb.PersistentClient(path=db_path)
 collection =
 client.get_or_create_collection(name=f'rag_experiment_{name}')
 DBS[name] = collection

=====
INDEXING PIPELINE
=====

def index_all_pipelines():
 """
 Index documents across all experimental configurations
 """

 for name, cfg in EXPERIMENTAL_PIPELINES.items():
 print(f"\n[EXPERIMENT] Indexing Pipeline: {name}")
 print(f" Model: {cfg['model']}")
 print(f" Chunk Size: {cfg['chunk_size']}")
 print(f" Strategy: {cfg['chunk_strategy']}")

 chunks, metas = load_and_chunk_folder(
 DOCS_FOLDER,
 cfg["chunk_size"],
 cfg["chunk_strategy"]
)

 if len(chunks) == 0:
 print(" ▲ No documents found")
 continue

 embeddings = EMBEDDERS[name].encode(chunks,
 show_progress_bar=True)
 embeddings = [e.tolist() for e in embeddings]

 ids = [f'{name}_{i}' for i in range(len(chunks))]

 try:
 DBS[name].add(
 ids=ids,
 documents=chunks,
 embeddings=embeddings,
 metadatas=metas,
)
 print(f" ✅ Indexed {len(chunks)} chunks")
 except Exception:
 print(f" ❌ Already indexed")
```

## PART 1: VECTOR DATABASE SETUP

```
python
```

```
DBS = {}
```

This is a dictionary that will store one vector database collection per experiment.

Think:

```
makefile
```

```
DBS = {
 "minILM_fixed100": <Chroma collection>,
 "E5small_fixed150": <Chroma collection>,
 "BGE_fixed256": <Chroma collection>,
}
```

### Why one database per pipeline?

Because:

- Each pipeline uses different embeddings
- You cannot mix embeddings from different models
- Otherwise similarity search becomes meaningless

★ Rule:

Vectors generated by model A must only be compared with vectors from model A.

### Looping over pipelines

```
python
```

```
for name in EXPERIMENTAL_PIPELINES.keys():
```

Example:

```
python
```

```
name = "minILM_fixed100"
```

## Database path

```
python

db_path = f"./chroma_db_{name}"
```

Creates a persistent folder on disk like:

```
bash

.chroma_db_minILM_fixed100/
.chroma_db_E5small_fixed150/
```

### ★ Why persistent?

- Data survives program restarts
- You don't re-index every time
- Production-ready behavior

## Create Chroma client

```
python

client = chromadb.PersistentClient(path=db_path)
```

This initializes ChromaDB, a vector database.

What Chroma stores:

- Embeddings (vectors)
- Original text chunks
- Metadata
- IDs

## Create or load a collection

```
python

collection = client.get_or_create_collection(
 name=f"rag_experiment_{name}"
)
```

A collection is like a table in SQL.

Each collection contains:

- vectors
- documents
- metadata

## Create or load a collection

```
python

collection = client.get_or_create_collection(
 name=f"rag_experiment_{name}"
)
```

A collection is like a table in SQL.

Each collection contains:

- vectors
- documents
- metadata

## Save collection

```
python

DBS[name] = collection
```

Now every pipeline has its own searchable vector store.

## PART 2: INDEXING PIPELINE (THIS IS WHERE KNOWLEDGE IS INGESTED)

python

```
def index_all_pipelines():
```

This function does **bulk knowledge ingestion**.

In plain English:

"Take documents → chunk them → embed them → store them in vector DB."

### Loop over each experimental pipeline

python

```
for name, cfg in EXPERIMENTAL_PIPELINES.items():
```

You are repeating the **same ingestion process** with:

- different embedding models
- different chunk sizes

This is **controlled experimentation**.

### Logging (important for experiments)

python

```
print(f"[EXPERIMENT] Indexing Pipeline: {name}")
```

This makes results reproducible and debuggable.

## Load and chunk documents

```
python

chunks, metas = load_and_chunk_folder(
 DOCS_FOLDER,
 cfg["chunk_size"],
 cfg["chunk_strategy"]
)
```

This gives you:

- chunks :

```
python
```

```
[
```

```
 "chunk text 1",
 "chunk text 2",
 ...
]
```

- metas :

```
python
```

```
[
```

```
 {"source": "paper1.pdf", "chunk": 0},
 {"source": "paper1.pdf", "chunk": 1},
]
```

★ This step transforms raw files into retrievable knowledge units.

## Empty document check

```
python

if len(chunks) == 0:
 print("⚠ No documents found")
 continue
```

Prevents:

- crashing
- indexing empty DBs
- misleading experiment results

## Create embeddings (CRITICAL STEP)

```
python
```

```
embeddings = EMBEDDERS[name].encode(chunks, show_progress_bar=True)
```

This converts:

```
arduino
```

```
"Neural networks are models..."
→ [0.034, -0.22, 0.91, ...]
```

★ This is where text becomes math.

## Convert to list

```
python
```

```
embeddings = [e.tolist() for e in embeddings]
```

Chroma expects **Python lists**, not NumPy arrays.

## Create unique IDs

```
python
```

```
ids = [f"{name}_{i}" for i in range(len(chunks))]
```

Example:

```
nginx
```

```
miniLM_fixed100_0
miniLM_fixed100_1
```

★ IDs ensure:

- no duplicates
- traceability
- safe updates

## Store everything in Chroma

```
python

DBS[name].add(
 ids=ids,
 documents=chunks,
 embeddings=embeddings,
 metadatas=metas,
)
```

This is THE MOMENT knowledge enters the RAG system.

Chroma now stores:

| Component  | Purpose            |
|------------|--------------------|
| embeddings | semantic search    |
| documents  | text sent to LLM   |
| metadatas  | source attribution |
| ids        | identity & updates |

## Why try / except ?

```
python

except Exception:
 print("⚠️ Already indexed")
```

- Prevents duplicate indexing
- Allows safe reruns
- Practical for development

# RETRIEVAL SYSTEM

```
=====
RETRIEVAL SYSTEM
=====

def retrieve(pipeline_name, query, top_k=TOP_K_DEFAULT):
 """
 Retrieve relevant chunks using specified pipeline
 """
 q_emb = EMBEDDERS[pipeline_name].encode([query])[0].tolist()

 result = DBS[pipeline_name].query(
 query_embeddings=[q_emb],
 n_results=top_k,
)

 docs = result["documents"][0]
 metas = result["metadatas"][0]

 return docs, metas
```

**BIG PICTURE (before code)**

Retrieval answers this question:  
"Given a user question, which parts of my documents are most relevant?"

The LLM does not search documents.  
The vector database does — using embeddings.

---

**WHAT THIS FUNCTION DOES (one sentence)**

Converts the user query into a vector, finds the *top-K most similar document chunks*, and returns them.

That's it.  
Everything else in RAG depends on this step.

---

**STEP-BY-STEP BREAKDOWN**

**1 Function signature**

python Copy code

```
def retrieve(pipeline_name, query, top_k=TOP_K_DEFAULT):
```

**Parameters**

| Parameter                  | Meaning                                      |
|----------------------------|----------------------------------------------|
| <code>pipeline_name</code> | Which experiment to use (MiniLM, E5, BGE...) |
| <code>query</code>         | User question                                |
| <code>top_k</code>         | How many chunks to retrieve                  |

★ Why `pipeline_name` matters:

- Each pipeline has its own embedding model
- Each pipeline has its own vector database
- Mixing them would break semantic meaning

## 2 Convert query into an embedding

python

 Copy code

```
q_emb = EMBEDDERS[pipeline_name].encode([query])[0].tolist()
```

### What's happening?

1. `encode([query])`
  - Converts the text query into a vector
  - Output shape: [1, embedding\_dim]
2. `[0]`
  - Extracts the single vector from the list
3. `.tolist()`
  - Converts NumPy array → Python list (required by Chroma)

### Conceptually:

csharp

 Copy code

```
"What is RAG?"
↓
[0.12, -0.44, 0.88, ...]
```

### ★ Critical rule of RAG

Query embeddings must be generated using the same model used for document embeddings.

### That's why we do:

python

 Copy code

```
EMBEDDERS[pipeline_name]
```

### 3 Query the vector database

python

```
result = DBS[pipeline_name].query(
 query_embeddings=[q_emb],
 n_results=top_k,
)
```

This tells Chroma:

"Find the `top_k` vectors in this collection that are *closest* to `q_emb`."

What does "closest" mean?

Chroma uses vector similarity (usually cosine similarity):

ini

```
similarity = cos(query_vector, document_vector)
```

- High similarity → similar meaning
- Low similarity → unrelated

👉 This is semantic search, not keyword matching.

### 4 Structure of result

Chroma returns a dictionary like:

python

```
{
 "documents": [[chunk1, chunk2, ...]],
 "metadatas": [[meta1, meta2, ...]],
 "ids": [[id1, id2, ...]],
 "distances": [[d1, d2, ...]]
}
```

Why double lists?

- Chroma supports batch queries
- You passed one query, so everything is inside `[0]`

## WHY `top_k` MATTERS

| <code>top_k</code> | Effect                         |
|--------------------|--------------------------------|
| Too small (1–2)    | Risk missing important context |
| Moderate (4–8)     | Balanced, recommended          |
| Too large (20+)    | Noisy context, LLM confusion   |

Your default of `5` is very reasonable.

### 5 Extract documents and metadata

```
python

docs = result["documents"][0]
metas = result["metadatas"][0]
```

Now you have:

```
python

docs = [
 "RAG is a technique that combines retrieval and generation...",
 "Vector databases store embeddings...",
]
```

```
python

metas = [
 {"source": "rag_paper.pdf", "chunk": 12},
 {"source": "ml_notes.txt", "chunk": 4},
]
```

📌 This is grounded knowledge.

### 6 Return results

```
python

return docs, metas
```

These are later sent to the LLM as context.

## EVALUATION

### Hit rate

```
def hit_rate(y_true, y_pred):
 ...
 Measures if ANY relevant document was retrieved
 ...
 hits = 0
 for true_docs, pred_docs in zip(y_true, y_pred):
 if any(doc in true_docs for doc in pred_docs):
 hits += 1
 return hits / len(y_true) if len(y_true) > 0 else 0
```

## 💡 What is Hit Rate? (Concept first)

Hit Rate answers ONE simple question:

❓ Did the retrieval system manage to retrieve at least one correct (relevant) document for a query?

It does NOT care about ranking, order, or how many were retrieved — only:

✓ Was there a hit or not?

## Why this matters in RAG

In RAG:

- If zero relevant chunks are retrieved →  
✗ LLM has no factual grounding → hallucination guaranteed
- If at least one relevant chunk is retrieved →  
✓ LLM has a chance to answer correctly

So hit rate is a minimum survival metric for RAG.

## 💡 Inputs to the function

python

 Copy code

```
y_true
y_pred
```

These are lists of lists.

Structure:

python

 Copy code

```
y_true = [
 [doc2, doc7], # relevant docs for query 1
 [doc5], # relevant docs for query 2
]

y_pred = [
 [doc9, doc7, doc1], # retrieved docs for query 1
 [doc3, doc4], # retrieved docs for query 2
]
```

## Expected Output

A number between 0 and 1

| Value | Meaning                                         |
|-------|-------------------------------------------------|
| 1.0   | Every query had $\geq 1$ relevant doc retrieved |
| 0.0   | No query retrieved a relevant doc               |
| 0.6   | 60% of queries had at least one hit             |

### What Hit Rate DOES and DOES NOT Measure

#### Measures:

- Minimum retrieval success
- Whether RAG is *even viable*

#### Does NOT measure:

- Ranking quality
- How many relevant docs
- Whether LLM uses them

### When is Hit Rate useful?

| Scenario                      | Usefulness                                                                                        |
|-------------------------------|---------------------------------------------------------------------------------------------------|
| Early RAG testing             |  Very useful |
| Comparing chunk sizes         |              |
| Comparing embedding models    |              |
| Fine-grained ranking analysis |              |

## Mean Reciprocal Rank (MRR)

```
def mean_reciprocal_rank(y_true, y_pred):
 """
 Measures ranking quality - how early relevant docs appear
 """

 reciprocal_ranks = []
 for true_docs, pred_docs in zip(y_true, y_pred):
 rr = 0
 for rank, doc in enumerate(pred_docs, start=1):
 if doc in true_docs:
 rr = 1 / rank
 break
 reciprocal_ranks.append(rr)
 return np.mean(reciprocal_ranks)
```

## 💡 What is Mean Reciprocal Rank (MRR)?

MRR answers this question:

❓ How early does the first correct (relevant) document appear in the retrieval list?

It cares about **ranking**, not just presence.

## ➡ Why “Reciprocal”?

If the first relevant document appears at:

| Rank      | Reciprocal Rank |
|-----------|-----------------|
| 1         | 1.0             |
| 2         | 0.5             |
| 3         | 0.33            |
| 5         | 0.2             |
| Not found | 0               |

Earlier = **higher score**

Later = **penalized**

## 📌 Why MRR matters in RAG

In RAG:

- LLMs **pay more attention** to earlier chunks
- Context windows are limited
- Top-1 or Top-3 chunks dominate the answer

So:

A relevant chunk at rank 1 is far better than rank 10

## 1 Function definition

python

```
def mean_reciprocal_rank(y_true, y_pred):
```

We compute average reciprocal rank across all queries.

## 2 Storage for per-query scores

python

```
reciprocal_ranks = []
```

Each query produces one score.

## 3 Loop over queries

python

```
for true_docs, pred_docs in zip(y_true, y_pred):
```

Process one query at a time.

## 4 Initialize reciprocal rank

python

```
rr = 0
```

Default = no relevant document found

## 5 Loop over retrieved docs with ranking

python

```
for rank, doc in enumerate(pred_docs, start=1):
```

- `rank` starts at 1 (not 0)
- `doc` is each retrieved chunk

## Concrete Example

### Input

```
python

y_true = [
 ["A"], # Query 1
 ["B"], # Query 2
 ["C"] # Query 3
]

y_pred = [
 ["A", "X"], # rank 1 → RR = 1.0
 ["X", "B"], # rank 2 → RR = 0.5
 ["X", "Y"] # not found → RR = 0
]
```

### Calculation

```
text

Query 1 → 1.0
Query 2 → 0.5
Query 3 → 0.0
```

```
python
```

```
MRR = (1.0 + 0.5 + 0.0) / 3 = 0.5
```

## 6 Check relevance

```
python

if doc in true_docs:
```

Is this retrieved document actually relevant?

## 7 Compute reciprocal rank

```
python

rr = 1 / rank
break
```

- First relevant document only
- Ignore later ones
- Break immediately

★ Key idea:

MRR only cares about the earliest correct hit.

## 8 Save this query's score

```
python

reciprocal_ranks.append(rr)
```

## 9 Average across queries

```
python

return np.mean(reciprocal_ranks)
```

This produces Mean Reciprocal Rank.

## Interpretation of MRR Values

| MRR  | Meaning              |
|------|----------------------|
| ~1.0 | Almost always rank-1 |
| 0.5  | Usually rank-2       |
| 0.25 | Often rank-4         |
| 0.0  | Retrieval failing    |

## How MRR differs from Hit Rate

| Metric   | Question                                       |
|----------|------------------------------------------------|
| Hit Rate | "Was <i>any</i> relevant doc retrieved?"       |
| MRR      | "How early was the <i>first</i> relevant doc?" |

Example:

- Retrieved at rank 10 → Hit Rate = 1, MRR = 0.1
- Retrieved at rank 1 → Hit Rate = 1, MRR = 1.0

Same hit rate — **very different quality.**

## precision recall at k

```
def precision_recall_at_k(y_true, y_pred, k=5):
 """
 Core retrieval metrics
 """
 recall_list = []
 precision_list = []

 for true_docs, pred_docs in zip(y_true, y_pred):
 top_k = pred_docs[:k]
 hits = len([doc for doc in top_k if doc in true_docs])

 recall_list.append(hits / len(true_docs) if len(true_docs) > 0 else 0)
 precision_list.append(hits / k if k > 0 else 0)

 return {
 f"Recall@{k}": np.mean(recall_list),
 f"Precision@{k}": np.mean(precision_list)
 }
```

## 🎯 What does `precision_recall_at_k` measure?

This function answers **two questions** for every query:

### 1. Recall@K

👉 *Did we retrieve most (or all) of the relevant chunks?*

### 2. Precision@K

👉 *How many of the retrieved chunks were actually relevant?*

## 🧠 Why this matters in RAG (big picture)

In RAG:

- LLM can ignore irrelevant chunks
- But it cannot invent missing information

So:

🔥 Recall is more important than precision in RAG

## 📌 Inputs (same pattern again)

python

```
y_true = [
 ["doc1", "doc2"], # relevant docs for query 1
 ["doc3"] # relevant docs for query 2
]

y_pred = [
 ["doc2", "docX", "docY", "doc1"], # retrieved docs for query 1
 ["docX", "doc3"] # retrieved docs for query 2
]
```

## ✳️ Line-by-Line Explanation

### 1 Function definition

python

```
def precision_recall_at_k(y_true, y_pred, k=5):
```

- Computes metrics only for the top `k` retrieved docs
- Default: `k = 5`

### 2 Storage lists

python

```
recall_list = []
precision_list = []
```

Each query gets:

- One recall score
- One precision score

### 3 Loop per query

python

```
for true_docs, pred_docs in zip(y_true, y_pred):
```

Process one query at a time.

#### 4 Take only top-K retrieved docs

```
python
top_k = pred_docs[:k]
```

Anything after rank `k` is ignored.

★ This mimics:

- LLM context limits
- Real retrieval behavior

#### 5 Count relevant hits

```
python
hits = len([doc for doc in top_k if doc in true_docs])
```

"How many of the top-K docs are actually relevant?"

#### 6 Compute Recall@K

```
python
recall_list.append(hits / len(true_docs) if len(true_docs) > 0 else 0)
```

Formula:

```
graphql
Recall@K = (# relevant docs retrieved in top K)

(total # of relevant docs)
```

Example:

- Relevant docs = 4
- Retrieved in top-5 = 3

$$\text{Recall@5} = 3 / 4 = 0.75$$

★ Recall answers:

"Did we retrieve everything important?"

## 7 Compute Precision@K

python

```
precision_list.append(hits / k if k > 0 else 0)
```

Formula:

mathematica

$$\text{Precision}@K = \frac{\text{\# relevant docs in top } K}{K}$$

Example:

- Relevant in top-5 = 2  
⇒ Precision@5 =  $2 / 5 = 0.4$

★ Precision answers:

"How clean were the retrieved results?"

## 8 Average across queries

python

```
return {
 f"Recall@{k}": np.mean(recall_list),
 f"Precision@{k}": np.mean(precision_list)
}
```

Final output:

python

```
{
 "Recall@5": 0.78,
 "Precision@5": 0.42
}
```

## 💡 Real RAG Scenario

You built a RAG system over Python documentation.

### User Query

"What does `zip()` do in Python?"

## 📚 Ground Truth (`y_true`)

These are the chunks that actually contain the correct answer.

```
python

y_true = [[
 "The zip() function returns an iterator of tuples",
 "zip() aggregates elements from multiple iterables",
 "zip() stops when the shortest iterable is exhausted"
]]
```

👉 There are 3 relevant chunks.

## 🔍 Retrieved Chunks (`y_pred`)

Your vector database returns this ranked list:

```
python

y_pred = [[
 "The zip() function returns an iterator of tuples",
 "map() applies a function to all items",
 "zip() aggregates elements from multiple iterables",
 "list comprehension syntax in Python",
 "zip() stops when the shortest iterable is exhausted"
]]
```

⚙ Let's run `precision_recall_at_k` with  $k = 3$

```
python
k = 3
top_k = y_pred[0][:3]
```

#### Top-3 Retrieved

```
text

1. zip() returns an iterator of tuples ✓
2. map() applies a function ✗
3. zip() aggregates elements from multiple iterables ✓
```

### 📋 Step-by-step Calculation

#### ◆ Hits

Relevant docs in top-3:

```
text

Hits = 2
```

#### ◆ Recall@3

```
text

Recall@3 = hits / total relevant
= 2 / 3
= 0.67
```

#### ★ Meaning:

We retrieved 67% of the required knowledge

#### ◆ Precision@3

```
text

Precision@3 = hits / k
= 2 / 3
= 0.67
```

#### ★ Meaning:

67% of retrieved chunks were useful



```
{
 "Recall@3": 0.67,
 "Precision@3": 0.67
}
```

---

---

## 🔥 Now let's improve retrieval (same example)

Increase `k = 5`

Top-5 retrieved:

text



1. `zip()` returns an iterator of tuples ✓
2. `map()` applies a function ✗
3. `zip()` aggregates elements from multiple iterables ✓
4. list comprehension syntax ✗
5. `zip()` stops when the shortest iterable is exhausted ✓

## New Metrics

text



Hits = 3  
 $\text{Recall@5} = 3 / 3 = 1.0$  ✓  
 $\text{Precision@5} = 3 / 5 = 0.6$

## 💡 Key Insight (THIS IS THE LESSON)

| Metric    | k=3  | k=5  |
|-----------|------|------|
| Recall    | 0.67 | 1.00 |
| Precision | 0.67 | 0.60 |

## What happened?

- Precision dropped slightly
  - Recall became perfect
- ↗️ For RAG this is GOOD

Why?

The LLM now sees *all necessary facts*  
Extra noise is manageable  
Missing facts cause hallucinations

## 💡 What the LLM Sees

### ✗ With Recall@3 = 0.67

LLM might miss:

“zip() stops at shortest iterable”

→ Partial or wrong explanation

---

### ✓ With Recall@5 = 1.0

LLM can answer fully:

What zip does  
How it aggregates  
When it stops

→ Correct, grounded answer

---

## ✳️ How your function captures this

python

```
hits = len([doc for doc in top_k if doc in true_docs])
```

- ✓ counts real knowledge
- ✓ ignores noise
- ✓ measures retrieval sufficiency

---

## ✍️ Why this example is “better”

- Realistic query
- Realistic chunks
- Shows why recall dominates precision
- Shows how top-K changes behavior
- Directly ties to hallucination risk

## nDCG

```
def ndcg_at_k(y_true, y_pred, k=5):
 """
 Normalized Discounted Cumulative Gain - ranking quality
 """
 ndcg_scores = []

 for true_docs, pred_docs in zip(y_true, y_pred):
 pred_docs_k = pred_docs[:k]

 # DCG
 dcg = sum([
 1 / np.log2(idx + 2) if doc in true_docs else 0
 for idx, doc in enumerate(pred_docs_k)
])

 # IDCG
 ideal_docs_k = true_docs[:k]
 idcg = sum([1 / np.log2(idx + 2) for idx, _ in
 enumerate(ideal_docs_k)])

 ndcg_scores.append(dcg / idcg if idcg > 0 else 0)

 return np.mean(ndcg_scores)
```



## Code Walkthrough (Line by Line)

### Function Signature

python

```
def ndcg_at_k(y_true, y_pred, k=5):
```

- `y_true` : relevant documents (ground truth)
- `y_pred` : retrieved documents (ranked)
- `k` : only evaluate top-k results

### Loop Over Queries

python

```
for true_docs, pred_docs in zip(y_true, y_pred):
```

Each query has:

- its **correct docs**
- its **retrieved ranking**

### Limit to Top-K

python

```
pred_docs_k = pred_docs[:k]
```

Only top-k matter — lower results won't influence the LLM much.

## 💡 What is nDCG (in plain English)?

nDCG measures how well your retrieval is ranked, not just *whether* relevant docs appear.

### Core idea:

- Relevant documents **earlier** in the list are **more valuable**
- Relevant documents **later** are still useful, but **discounted**
- Score is **normalized to [0, 1]**
  - **1.0** = perfect ranking
  - **0.0** = useless ranking

This is **critical** in RAG because:

- The LLM reads **top chunks first**
- Early facts dominate the answer

## 🔑 Two Core Concepts

### 1 DCG — Discounted Cumulative Gain

"How much useful information did we retrieve, considering rank?"

### 2 IDCG — Ideal DCG

"What would the score be if ranking were perfect?"

### 3 nDCG = DCG / IDCG

"How close are we to ideal?"

## 📌 Code Walkthrough (Line by Line)

### Function Signature

python

```
def ndcg_at_k(y_true, y_pred, k=5):
```

- `y_true` : relevant documents (ground truth)
- `y_pred` : retrieved documents (ranked)
- `k` : only evaluate top-k results

### Loop Over Queries

python

```
for true_docs, pred_docs in zip(y_true, y_pred):
```

Each query has:

- its **correct docs**
- its **retrieved ranking**

### Limit to Top-K

python

```
pred_docs_k = pred_docs[:k]
```

Only top-k matter — lower results won't influence the LLM much.

## DCG Calculation

python

```
dcg = sum([
 1 / np.log2(idx + 2) if doc in true_docs else 0
 for idx, doc in enumerate(pred_docs_k)
])
```



What this does:

| Rank | idx | Weight                       |
|------|-----|------------------------------|
| 1    | 0   | $1 / \log_2(2) = 1.0$        |
| 2    | 1   | $1 / \log_2(3) \approx 0.63$ |
| 3    | 2   | $1 / \log_2(4) = 0.5$        |
| 4    | 3   | $1 / \log_2(5) \approx 0.43$ |

❖ Earlier ranks matter more

If document is:

- relevant → add discounted gain
- irrelevant → add 0

## IDCG Calculation (Ideal Ranking)

python

```
ideal_docs_k = true_docs[:k]
idcg = sum([1 / np.log2(idx + 2) for idx, _ in enumerate(ideal_docs_k)])
```

This assumes:

- All relevant documents appear
  - In the **best possible order**
- ↗ This is your **maximum achievable DCG**

---

## Normalize

python

```
ndcg_scores.append(dcg / idcg if idcg > 0 else 0)
```

- Makes score **comparable across queries**
- Handles different **numbers of relevant docs**

---

## Final Score

python

```
return np.mean(ndcg_scores)
```

Average across all queries.

## 🔥 Concrete Numerical Example (RAG-style)

### Query

```
"What does zip() do in Python?"
```

### Ground Truth

```
python

true_docs = [
 "zip returns tuples",
 "zip aggregates iterables",
 "zip stops at shortest iterable"
]
```

### Retrieved Ranking

```
python

pred_docs = [
 "zip aggregates iterables", # ✅ rank 1
 "map applies function", # ❌ rank 2
 "zip stops at shortest iterable",# ✅ rank 3
 "list comprehensions", # ❌ rank 4
 "zip returns tuples" # ✅ rank 5
]
```

### Calculate DCG@5

| Rank | Relevant? | Score                         |
|------|-----------|-------------------------------|
| 1    | ✅         | $1 / \log_2(2) = 1.0$         |
| 2    | ❌         | 0                             |
| 3    | ✅         | $1 / \log_2(4) = 0.5$         |
| 4    | ❌         | 0                             |
| 5    | ✅         | $1 / \log_2(6) \approx 0.386$ |

```
text
```

```
DCG = 1.0 + 0 + 0.5 + 0 + 0.386 = 1.886
```

## Calculate IDCG@5 (Perfect Ranking)

| Rank | Score |
|------|-------|
| 1    | 1.0   |
| 2    | 0.63  |
| 3    | 0.5   |

text

IDCG  $\approx$  2.13

## Final nDCG

text

nDCG = 1.886 / 2.13  $\approx$  0.88

⌚ Very good ranking, but not perfect.

## 💡 Why nDCG is CRITICAL in RAG

| Metric    | What it tells                             |
|-----------|-------------------------------------------|
| Recall    | Did we retrieve the info?                 |
| Precision | How noisy is retrieval?                   |
| nDCG      | Did we rank info correctly for LLM usage? |

**High Recall + Low nDCG = ✗**

- Relevant docs buried
- LLM may miss them

**High Recall + High nDCG = ✓**

- Facts appear early
- LLM grounded & confident

# Summary of metrics:

**1 MRR — Mean Reciprocal Rank**

❓ What does it measure?  
How early does the first correct document appear?

---

💡 Intuition  
"Did I find the right answer quickly?"  
If the first relevant chunk is:

- Rank 1 → perfect
- Rank 5 → worse
- Not found → zero

---

💡 Formula (conceptual)  
ini  
`MRR = average(1 / rank_of_first_relevant_doc)`

---

📌 Example  
Retrieved list:  
markdown  
1. ✗ irrelevant  
2. ✗ irrelevant  
3. ✓ relevant  
Reciprocal rank = 1 / 3 = 0.33

---

💡 Why MRR matters in RAG

- LLMs pay more attention to earlier chunks
- Earlier = stronger influence on answer
- Late relevant docs = weaker grounding

✓ High MRR → low hallucination risk  
✗ Low MRR → relevant info buried too deep

## 2 Precision@K — How clean is retrieval?

### ❓ What does it measure?

What fraction of retrieved chunks are actually relevant?

### 💡 Intuition

"How much noise did I retrieve?"

### ⚠️ Formula

mathematica

```
Precision@K = relevant_docs_in_top_K / K
```

### 📌 Example (K=5)

Retrieved:

3 relevant, 2 irrelevant

Precision@5 = 3 / 5 = 0.6

### 🔍 Why Precision matters in RAG

- High precision → less distracting context
- Low precision → LLM sees noise

But ⚠️ precision alone is dangerous:

- You can be precise but miss important facts

📌 In RAG:

Recall > Precision, but precision still matters

### 3 nDCG@K (Normalized Discounted Cumulative Gain)

❓ Question it answers

| Did I put the relevant documents in the best possible order?

#### Important idea

- Relevant docs earlier = more valuable
- Relevant docs later = less useful
- Uses log to reduce score smoothly

#### Step-by-step (K = 5)

Actual ranking:

nginx

Copy

Rank 2 → relevant  
Rank 4 → relevant

Their contributions:

matlab

Copy

Rank 2 →  $1 / \log_2(2 + 1) = 0.63$   
Rank 4 →  $1 / \log_2(4 + 1) = 0.43$

DCG ≈ 1.06

Ideal ranking (best possible):

nginx

Copy

Rank 1 → relevant  
Rank 2 → relevant

IDCG:

matlab

Copy

$1 / \log_2(2) + 1 / \log_2(3) = 1 + 0.63 = 1.63$

Final score

ini

Copy

$nDCG = DCG / IDCG = 1.06 / 1.63 = 0.65$

## 🔥 One-table summary (very important)

| Metric      | What it checks          | Simple meaning                   |
|-------------|-------------------------|----------------------------------|
| MRR         | First relevant position | "Did I find a good doc early?"   |
| Precision@K | Relevant fraction       | "How much noise did I retrieve?" |
| nDCG@K      | Ranking quality         | "How well did I order results?"  |

## 🎯 Why all three are needed (RAG view)

- MRR high → critical info appears early
- Precision high → less hallucination
- nDCG high → optimal chunk ordering

👉 A retriever can:

- Have **good precision** but **bad MRR**
- Have **good MRR** but **bad nDCG**

That's why you measure **all three**.

# Retrieval–Generation Coupling

```

def context_utilization(responses, retrieved_docs, top_k=5):
 """
 CRITICAL METRIC: How much of the retrieved context is used in the
 answer?
 Measures retrieval → generation coupling
 """
 scores = []

 for resp, docs in zip(responses, retrieved_docs):
 resp_words = set(nltk.word_tokenize(resp.lower()))
 doc_words = set(
 word for d in docs[:top_k]
 for word in nltk.word_tokenize(d.lower())
)

 overlap = len(resp_words & doc_words) / len(resp_words) if
len(resp_words) > 0 else 0
 scores.append(overlap)

 return np.mean(scores)

def answer_relevance(responses, queries):
 """
 Does the answer actually address the query?
 """
 scores = []

 for resp, query in zip(responses, queries):
 resp_words = set(nltk.word_tokenize(resp.lower()))
 query_words = set(nltk.word_tokenize(query.lower()))

 overlap = len(resp_words & query_words)
 scores.append(overlap / len(query_words) if len(query_words) > 0
else 0)

 return np.mean(scores)

def retrieval_generation_correlation(retrieval_scores,
generation_scores):
 """
 META-METRIC: Do better retrievals lead to better answers?
 This is the KEY insight of your project
 """
 if len(retrieval_scores) < 2 or len(generation_scores) < 2:
 return 0.0

 correlation = np.corrcoef(retrieval_scores, generation_scores)[0, 1]
 return correlation if not np.isnan(correlation) else 0.0

```

## What is “Retrieval–Generation Coupling”?

In a RAG system there are **two stages**:

1. **Retrieval** → get chunks from the vector DB
2. **Generation** → LLM writes an answer using those chunks

💡 **Retrieval–generation coupling asks:**

“Did the generator actually USE what the retriever fetched?”

A system can:

- retrieve good chunks ✗ but ignore them
- retrieve bad chunks ✗ and hallucinate
- retrieve good chunks ✓ and use them properly ✓

This section measures **that relationship**.

## 1 context\_utilization

python

 Copy code

```
def context_utilization(responses, retrieved_docs, top_k=5):
```

### What it measures (plain English)

| How much of the generated answer actually comes from the retrieved documents?

If the answer uses words that appear in retrieved chunks, coupling is strong.

If it invents new words/ideas, coupling is weak.

### Step-by-step with an example

#### Retrieved chunks (top-2)

nginx

 Copy code

Chunk 1: "Transformers use self-attention"

Chunk 2: "Self-attention captures global context"

#### Generated answer

arduino

 Copy code

"Transformers rely on self-attention to capture global context"

## Tokenization

### Answer words

```
php
```

```
{transformers, rely, self-attention, capture, global, context}
```

### Retrieved words

```
python
```

```
{transformers, use, self-attention, captures, global, context}
```

## Overlap

```
php
```

```
Common words = {transformers, self-attention, global, context}
```

### Overlap score:

```
4 / 6 = 0.67
```

## Code logic explained

python

```
resp_words = set(nltk.word_tokenize(resp.lower()))
```

→ unique words in the answer

python

```
doc_words = set(
 word for d in docs[:top_k]
 for word in nltk.word_tokenize(d.lower())
)
```

→ unique words from retrieved chunks

python

```
overlap = len(resp_words & doc_words) / len(resp_words)
```

fraction of answer grounded in retrieved content

## Interpretation

| Score                                | Meaning            |
|--------------------------------------|--------------------|
| High ( $\approx 0.6\text{--}0.9$ )   | Answer is grounded |
| Medium ( $\approx 0.3\text{--}0.5$ ) | Partial grounding  |
| Low ( $<0.2$ )                       | Hallucination risk |

## 2 answer\_relevance

python

 Copy code

```
def answer_relevance(responses, queries):
```

### What it measures

Does the answer actually talk about what the user asked?

This is NOT about documents — it's about query ↔ answer alignment.

### Example

#### Query

arduino

 Copy code

```
"What is self-attention?"
```

#### Answer A (good)

css

 Copy code

```
"Self-attention allows a model to focus on different parts of a sentence"
```

#### Answer B (bad)

arduino

 Copy code

```
"Transformers were introduced in 2017"
```

## Token overlap

Query words:

```
lua

{what, self-attention}
```

Answer A overlap:

```
lua

{self-attention} → 1/2 = 0.5
```

Answer B overlap:

```
{ } → 0
```

## Code logic

```
python

resp_words = set(...)
query_words = set(...)
```

```
python

overlap = len(resp_words & query_words)
score = overlap / len(query_words)
```

## Interpretation

| Score | Meaning                 |
|-------|-------------------------|
| High  | Answer is on-topic      |
| Low   | Answer drifts           |
| Zero  | Answer ignores question |

### 3 retrieval\_generation\_correlation

python

```
def retrieval_generation_correlation(retrieval_scores, generation_scores):
```

#### What this REALLY measures (important)

- When retrieval quality improves, does answer quality also improve?

This is a meta-metric — it evaluates your entire RAG design, not a single answer.

#### Example

##### Pipeline A

| Query | Retrieval Score | Generation Score |
|-------|-----------------|------------------|
| Q1    | 0.9             | 0.85             |
| Q2    | 0.8             | 0.78             |
| Q3    | 0.3             | 0.25             |

- Strong positive correlation
- Retrieval directly influences generation

##### Pipeline B

| Query | Retrieval Score | Generation Score |
|-------|-----------------|------------------|
| Q1    | 0.9             | 0.2              |
| Q2    | 0.3             | 0.9              |
| Q3    | 0.7             | 0.4              |

- Weak or negative correlation
- LLM ignores retrieval

#### Code logic

python

```
correlation = np.corrcoef(retrieval_scores, generation_scores)[0, 1]
```

- +1 → strong positive relationship
- 0 → no relationship
- 1 → inverse relationship



| 🔥 Summary (very important)       |                                         |
|----------------------------------|-----------------------------------------|
| Metric                           | What it checks                          |
| Context Utilization              | Did the answer use retrieved chunks?    |
| Answer Relevance                 | Did the answer address the question?    |
| Retrieval-Generation Correlation | Does better retrieval → better answers? |

## Hallucination vs Coverage

### What is “Hallucination vs Coverage” in RAG?

When an LLM answers using retrieved documents, **two things can go wrong**:

1. ✗ **Hallucination**  
→ The model says things **not present** in retrieved documents.
2. ✗ **Under-coverage**  
→ The model retrieves useful info but **fails to include** it in the answer.

So we care about **two opposing goals**:

- **Groundedness** → Don't invent facts
- **Coverage** → Don't ignore retrieved info

This section measures **that balance**.

### 1 groundedness

```
python
def groundedness(responses, retrieved_docs, top_k=5):
```

#### What it measures (plain English)

What fraction of the answer is supported by retrieved documents?

High groundedness =

- ✓ low hallucination risk
- ✓ answer is faithful to sources

## Simple example

### Retrieved docs

```
vbnet
```

```
Doc: "Transformers use self-attention for sequence modeling"
```

### Answer A (grounded)

```
arduino
```

```
"Transformers use self-attention for modeling sequences"
```

Overlap is high → grounded

### Answer B (hallucinated)

```
arduino
```

```
"Transformers were invented in 2022 and use quantum layers"
```

Overlap is very low → hallucination

## Code logic

```
python
```

```
resp_words = set(nltk.word_tokenize(resp.lower()))
```

→ unique words in the answer

```
python
```

```
doc_words = set(
 word for d in docs[:top_k]
 for word in nltk.word_tokenize(d.lower())
)
```

→ unique words in retrieved chunks

```
python
```

```
overlap = len(resp_words & doc_words) / len(resp_words)
```

fraction of answer words supported by retrieval

## 2 hallucination\_rate

python

```
def hallucination_rate(responses, retrieved_docs, top_k=5):
```

### What it measures

What fraction of the answer is NOT supported by retrieved docs?

This is the direct opposite of groundedness.

### Example

Answer words:

yaml

```
{transformers, quantum, invented, 2022}
```

Retrieved words:

lua

```
{transformers, self-attention, modeling}
```

Unsupported words:

yaml

```
{quantum, invented, 2022}
```

Hallucination rate:

3 / 4 = 0.75

⚠ Very high hallucination risk.

### Code logic

python

```
unsupported = len(resp_words - doc_words)
```

This means:

words in answer minus words in docs

python

```
unsupported / len(resp_words)
```



### 3 coverage\_score

python

```
def coverage_score(responses, retrieved_docs, top_k=5):
```

#### What it measures

How much of the retrieved information actually appears in the answer?

This is the other side of the coin.

#### Why this matters

Imagine retrieval returns great chunks, but the answer only uses one sentence from them.

- ✓ No hallucination
- ✗ But missing important details

That's low coverage.

#### Example

##### Retrieved docs

lua

```
{self-attention, transformers, positional-encoding, multi-head, scaling}
```

##### Answer

arduino

```
"Transformers use self-attention"
```

##### Overlap:

lua

```
{transformers, self-attention}
```

##### Coverage:

```
2 / 5 = 0.4
```

Answer is safe but incomplete.



## Code logic

python

```
coverage = len(resp_words & doc_words) / len(doc_words)
```

➡ fraction of retrieved info that made it into the answer

## Interpretation

| Coverage         | Meaning          |
|------------------|------------------|
| High (>0.6)      | Complete         |
| Medium (0.3–0.5) | Partial          |
| Low (<0.2)       | Wasted retrieval |

## 🔥 The CORE trade-off (very important)

| Metric                          | Risk                      |
|---------------------------------|---------------------------|
| High groundedness, low coverage | Safe but shallow          |
| High coverage, low groundedness | Detailed but hallucinated |
| High both                       | Ideal RAG system          |

👉 Best RAG systems maximize BOTH groundedness and coverage.

## Why these metrics matter for your project

Most RAG papers say:

“Our system reduces hallucination”

You can prove it quantitatively.

You’re not evaluating:

- just retrieval ✗
- just generation ✗

You’re evaluating faithfulness of the entire pipeline.

## 4. GENERATION QUALITY METRICS

```
def fluency_score(responses):
 """
 Readability and sentence structure quality
 """

 readability_scores = []
 coherence_scores = []

 for resp in responses:
 try:
 readability_scores.append(flesch_reading_ease(resp))
 except:
 readability_scores.append(0)

 sentences = nltk.sent_tokenize(resp)
 words = nltk.word_tokenize(resp)
 coherence = len(words) / len(sentences) if len(sentences) > 0 else
0
 coherence_scores.append(coherence)

 return {
 "Readability (Flesch)": np.mean(readability_scores),
 "Coherence (words/sentence)": np.mean(coherence_scores)
 }

def answer_length_stats(responses):
 """
 Simple length metrics
 """

 lengths = [len(resp.split()) for resp in responses]
 return {
 "Avg Answer Length": np.mean(lengths),
 "Min Length": np.min(lengths),
 "Max Length": np.max(lengths)
 }
```

## Why do we need *Generation Quality* metrics?

Even if:

- retrieval is perfect ✓
- groundedness is high ✓

the answer can still be:

- hard to read ✗
- rambling ✗
- too short or too long ✗

So we evaluate how good the language itself is, independent of correctness.

### 1 fluency\_score(responses)

python

```
def fluency_score(responses):
```

This measures how readable and coherent the answers are.

It has two sub-metrics:

1. Readability
2. Coherence

### A) Readability (Flesch Reading Ease)

python

```
readability_scores.append(flesch_reading_ease(resp))
```

What is Flesch Reading Ease?

It's a classical NLP readability metric.

| Score  | Meaning           |
|--------|-------------------|
| 90-100 | Very easy (kids)  |
| 60-70  | Plain English     |
| 30-50  | Academic          |
| < 30   | Very hard to read |

## Example

### Easy

css

Transformers use self-attention to process text.

→ High Flesch score

### Hard

nginx

Transformational architectures employing multi-headed attention mechanisms...

→ Lower score

## Why use it in RAG?

- Academic RAG → moderate readability is OK
- User-facing chatbot → high readability preferred

This lets you quantify clarity, not just correctness.

## Why try/except ?

python

```
try:
 flesch_reading_ease(resp)
except:
 readability_scores.append(0)
```

- Some answers may be too short or malformed
- Prevents evaluation from crashing

## B) Coherence (words per sentence)

```
python

sentences = nltk.sent_tokenize(resp)
words = nltk.word_tokenize(resp)
coherence = len(words) / len(sentences)
```

Copy code

What does this measure?

Average sentence length

This is a proxy for coherence.

### Interpretation

| Words / Sentence | Meaning  |
|------------------|----------|
| 8–15             | Clear    |
| 15–25            | Academic |
| >30              | Rambling |
| <5               | Choppy   |

### Example

X Poor coherence:

```
nginx

Transformers exist. Attention. Architecture. Model.
```

Copy code

X Poor coherence:

```
python

Transformers use self-attention mechanisms in a deeply hierarchical layered architecture enabling
```

Copy code

✓ Good coherence:

```
css

Transformers use self-attention to model long-range dependencies in text.
```

Copy code

### Final output of fluency\_score

```
python

{
 "Readability (Flesch)": 52.3,
 "Coherence (words/sentence)": 18.7
}
```

This tells you:

- Is the answer readable?
- Is it structured well?

### 2 answer\_length\_stats(responses)

```
python

def answer_length_stats(responses):

This evaluates verbosity control.
```

#### Why length matters in RAG

| Problem      | Example               |
|--------------|-----------------------|
| Too short    | Misses important info |
| Too long     | Hallucination risk    |
| Inconsistent | Bad UX                |

#### Code logic

```
python

lengths = [len(resp.split()) for resp in responses]

Counts words in each answer.

python

return {
 "Avg Answer Length": np.mean(lengths),
 "Min Length": np.min(lengths),
 "Max Length": np.max(lengths)
}
```

### Example output

json

```
{
 "Avg Answer Length": 182,
 "Min Length": 63,
 "Max Length": 418
}
```

### Interpretation

- Large spread → unstable generation
- Very long max → possible hallucination
- Very short min → retrieval not used well

### How this fits into the FULL evaluation

| Layer              | What it checks               |
|--------------------|------------------------------|
| Retrieval          | Did we fetch correct chunks? |
| Coupling           | Did the answer use them?     |
| Faithfulness       | Did it hallucinate?          |
| Generation quality | Is it readable & well-sized? |